Data Engineering with Apache Spark

Markus Dale, medale@asymmetrik.com

May 2019

Intro, Slides And Code

- Slides: https://github.com/medale/
- Scala Spark Code Examples: https://github.com/medale/

Data engineering

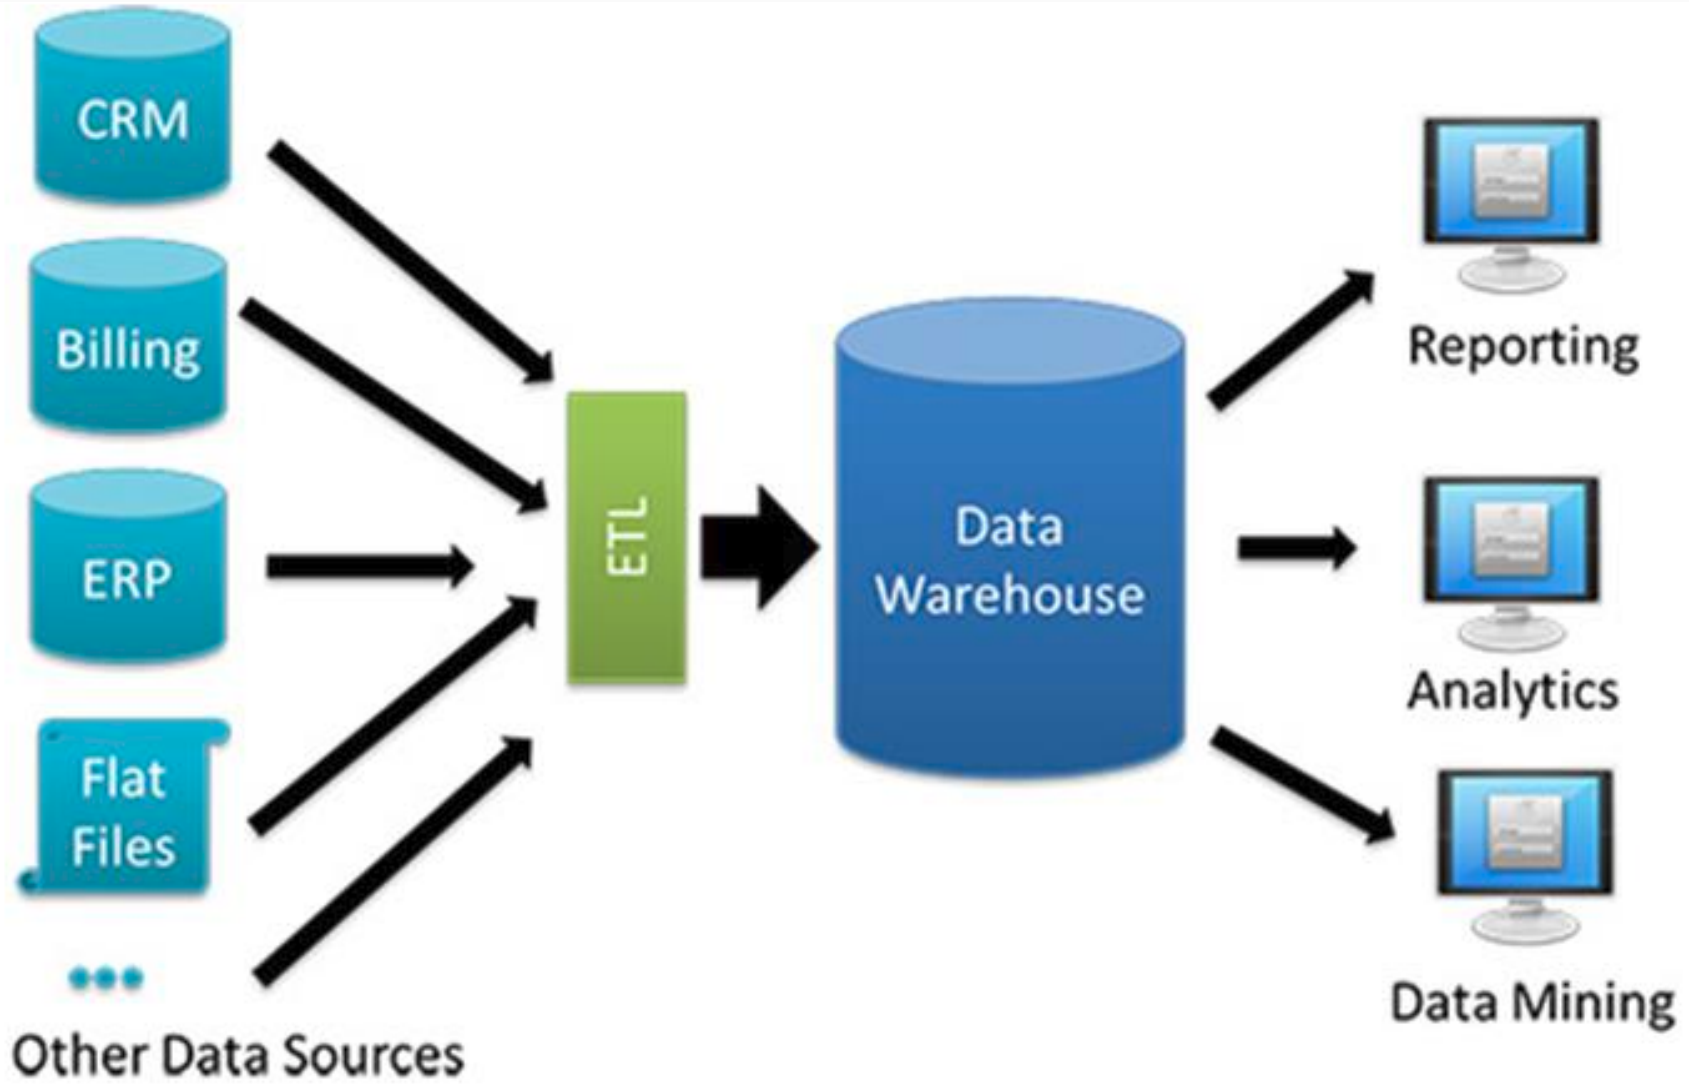

Data engineering on small dataset

Figure 1: Laptop

Data engineering for larger dataset (Vertical Scaling)

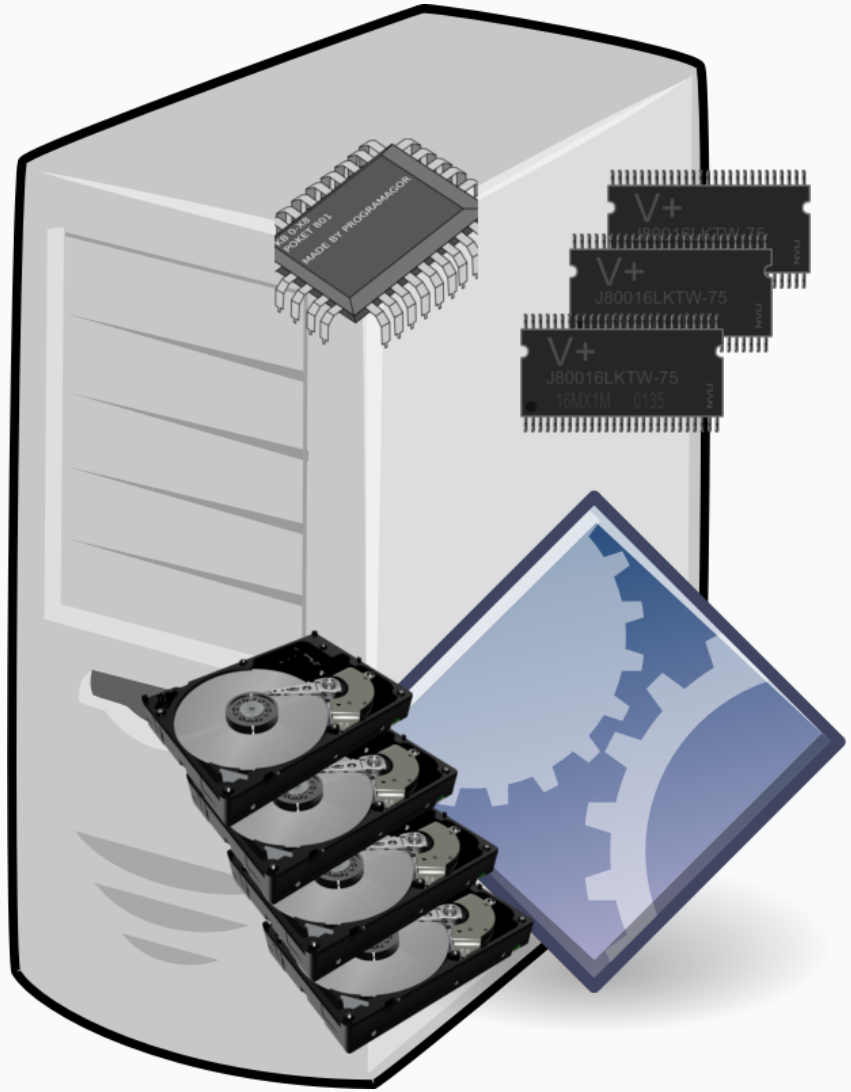

Figure 2: Beefed-up Server

Data engineering for large datasets (Horizontal Scaling)

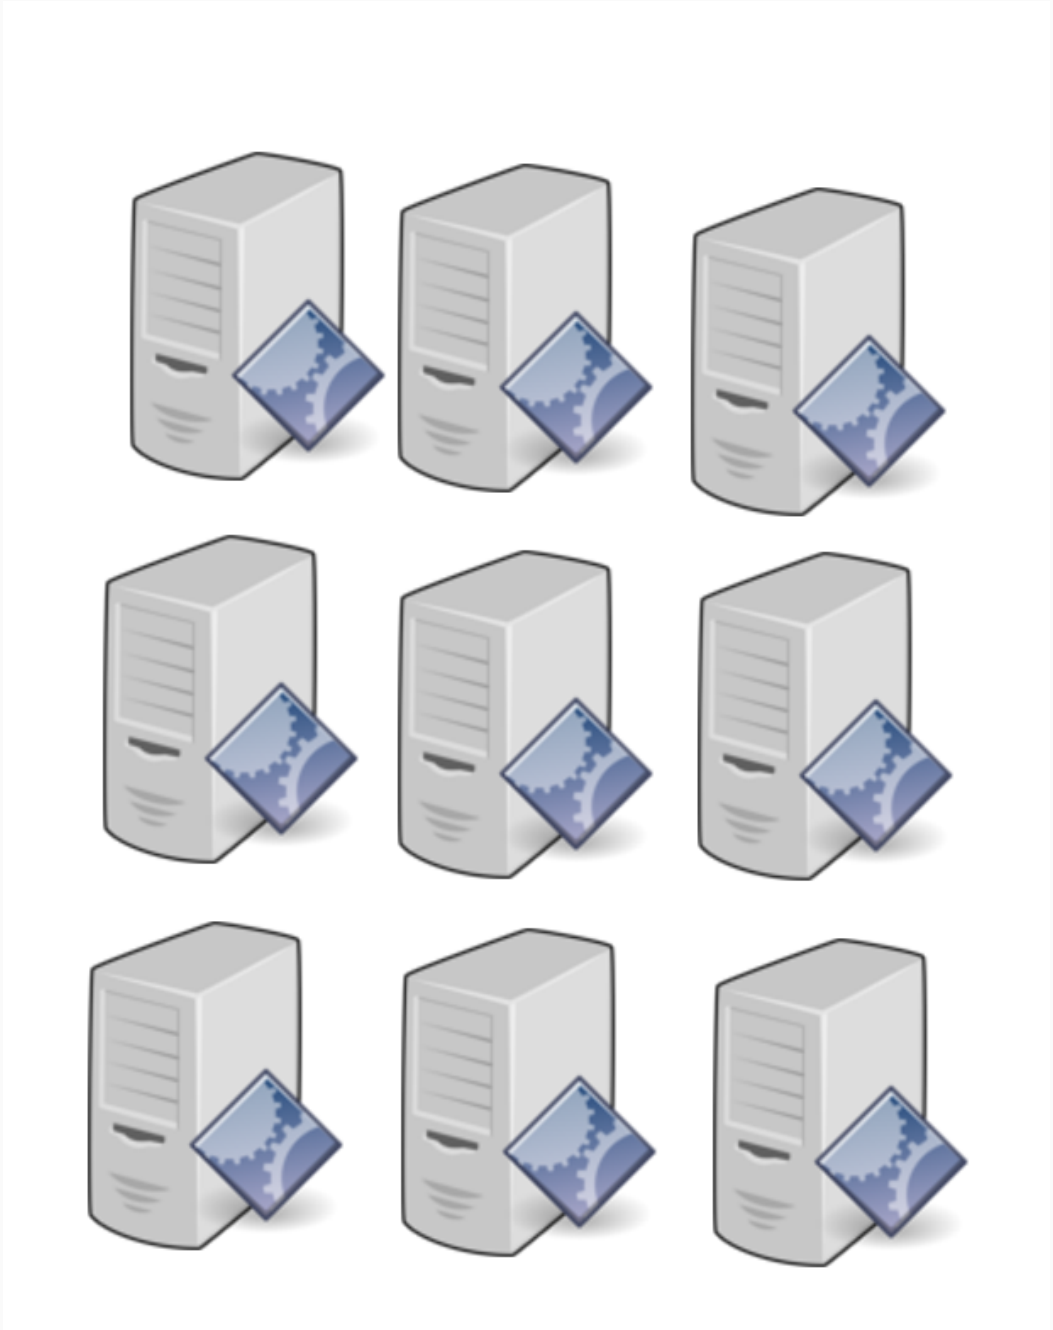

Cluster Manager

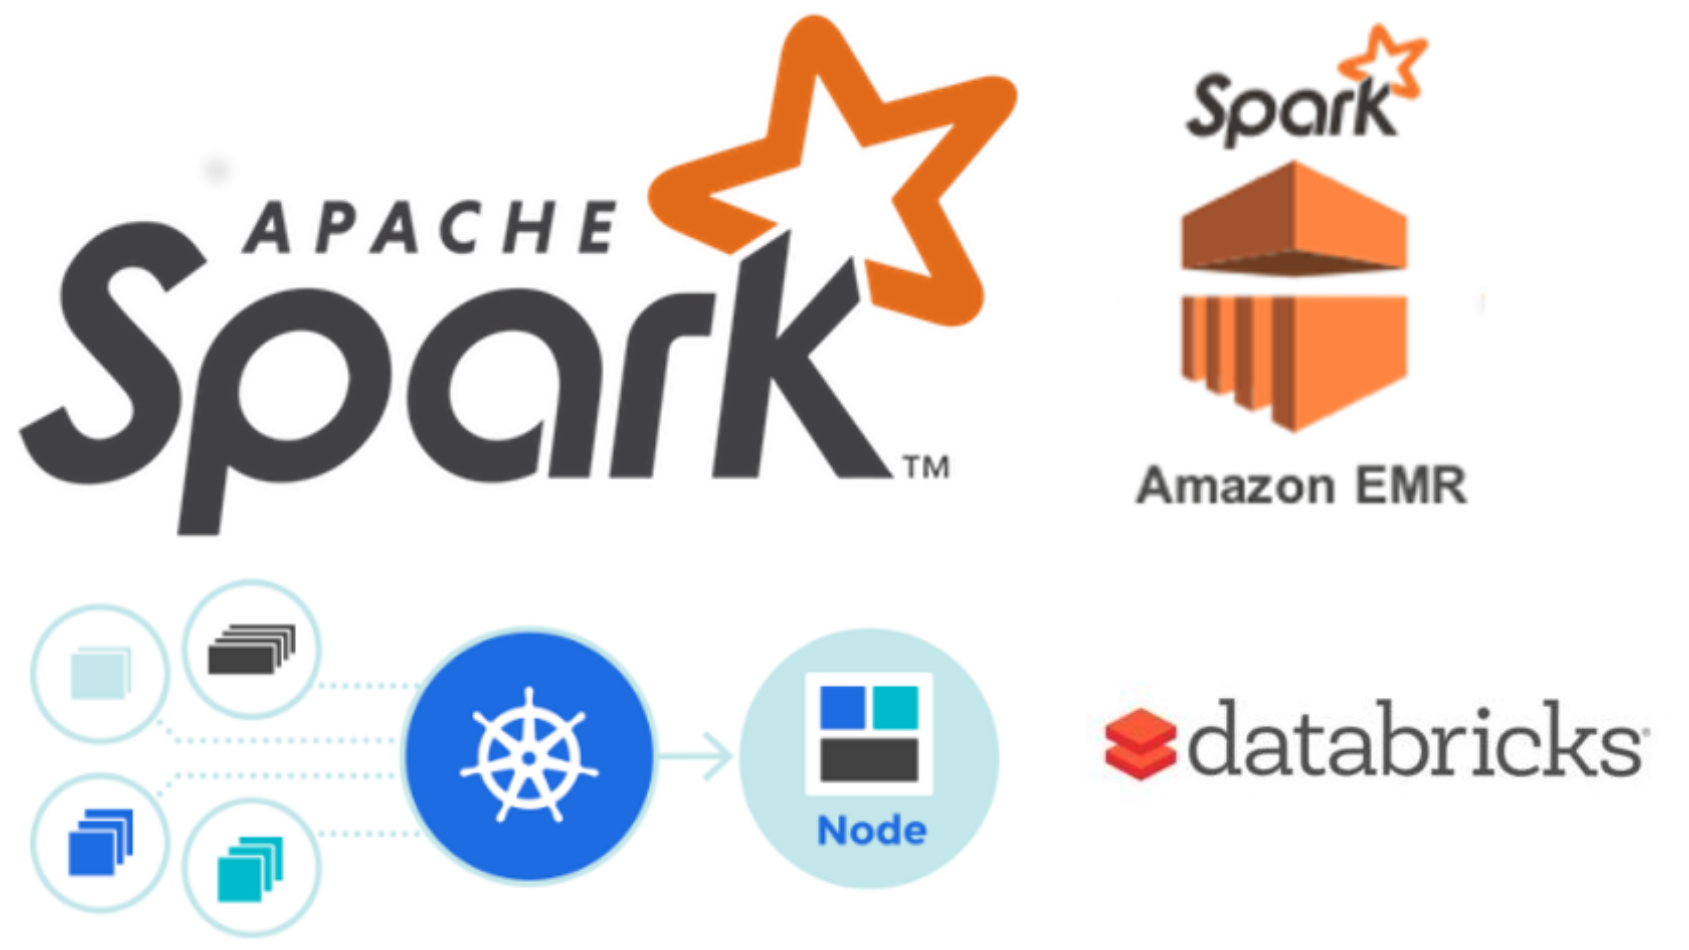

Resilient Distributed Datasets (RDDs)

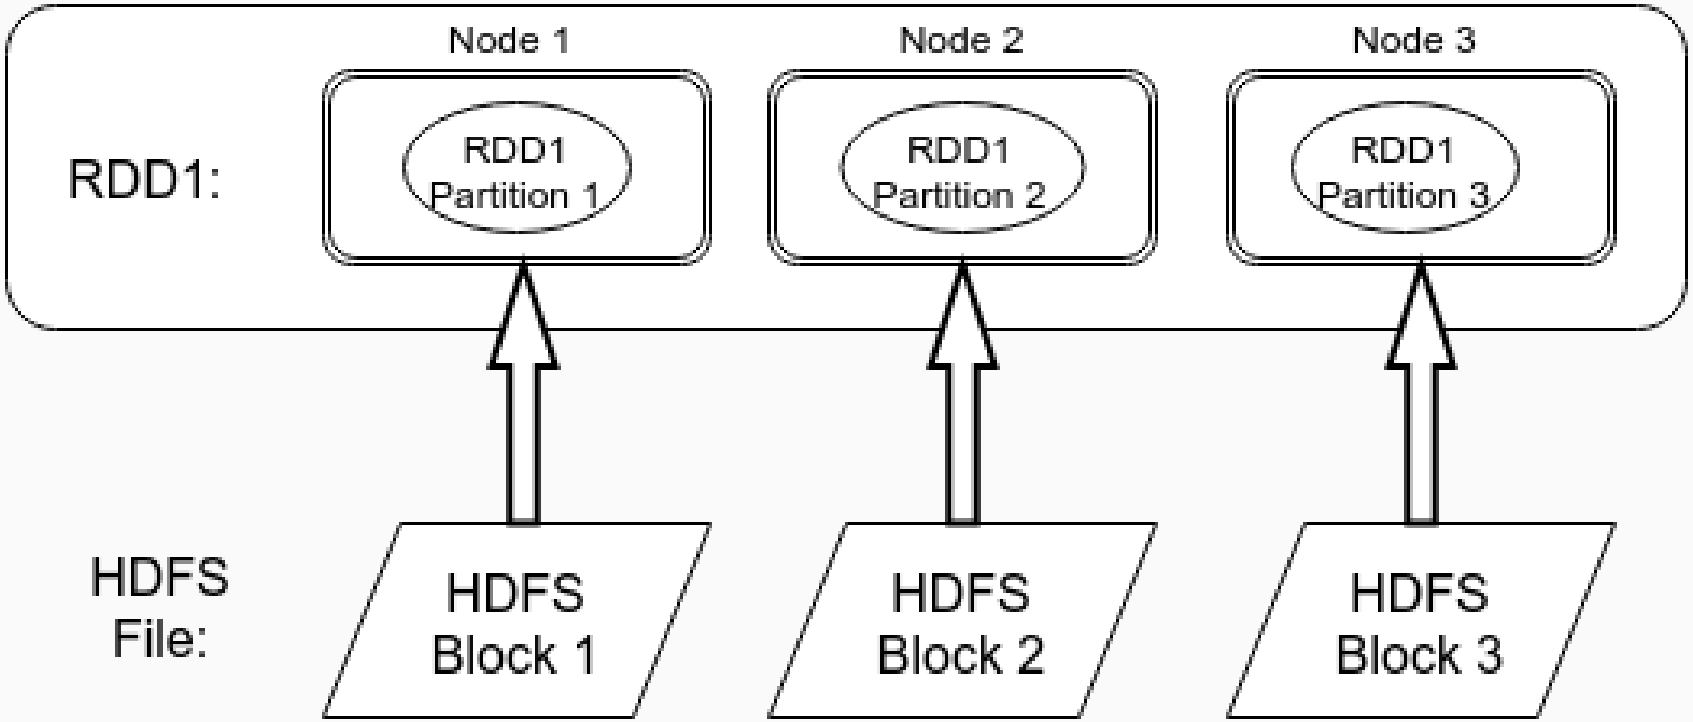

Anatomy of a Spark Application

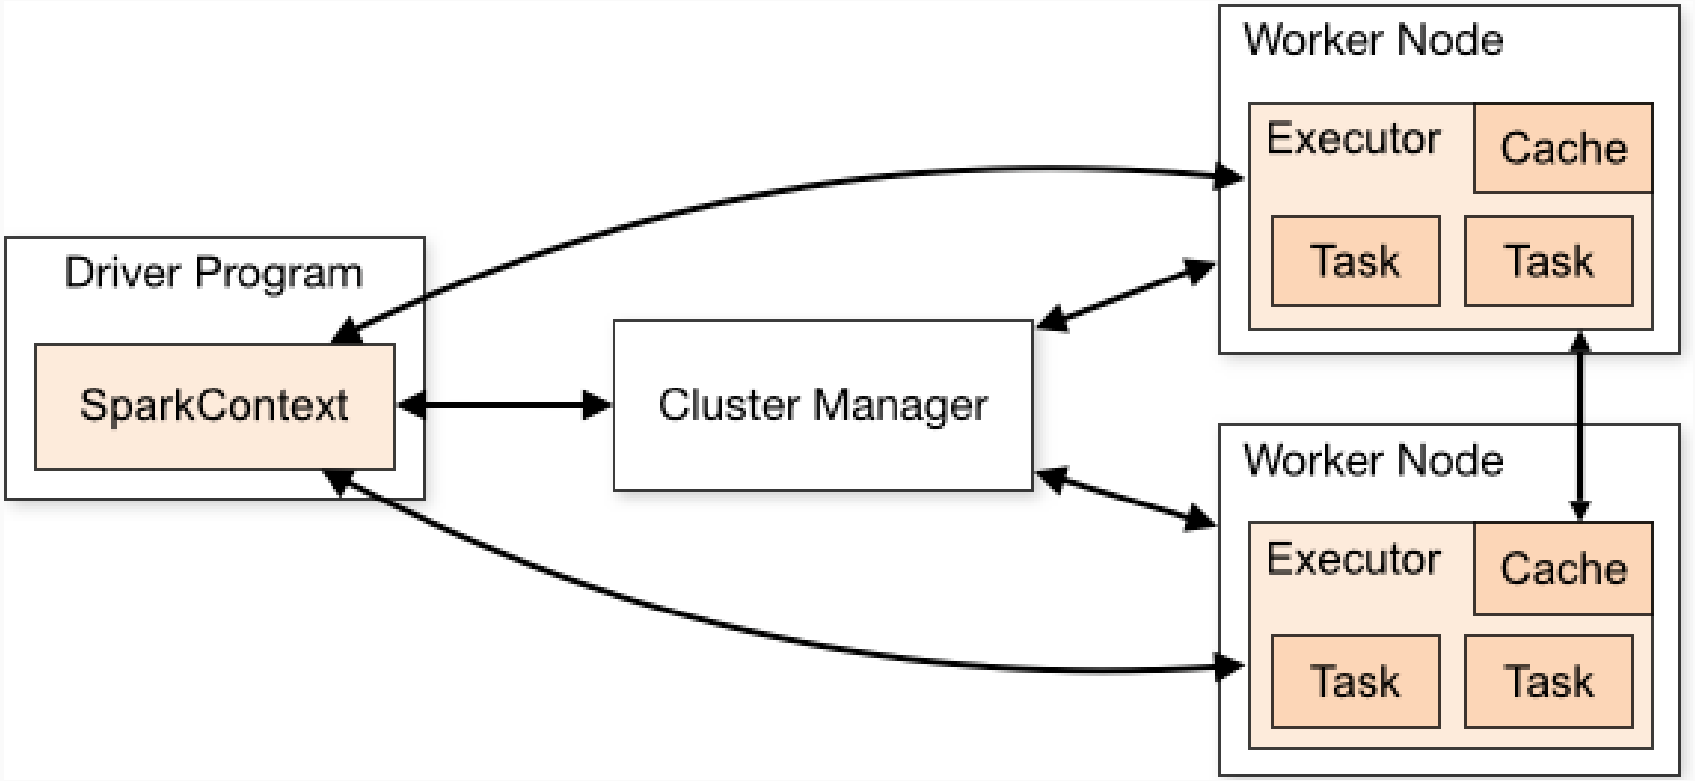

Source: Apache Spark website

Hello, Spark World!

```scala
import org.apache.spark.sql.SparkSession

object HelloSparkWorld {

  def process(spark: SparkSession): (Long,Long) = {
    val records = spark.read.json( path = "file:///datasets/github/data")
    records.cache()
    val totalEventCount = records.count()

    val prs = records.where(records("type") === "PullRequestEvent")
    val pullRequestEventCount = prs.count()

    records.unpersist()
    (totalEventCount, pullRequestEventCount)
  }

  def main(args: Array[String]): Unit = {
    val spark = SparkSession.builder().
      appName( name = "HelloSparkWorld").
      getOrCreate()
    process(spark)
  }
}
```

Starting Spark Standalone Cluster Manager

```
# Start on master
$SPARK_HOME/sbin/start-master.sh --host 192.168.1.230

# Start one or more workers
$SPARK_HOME/sbin/start-slave.sh spark://192.168.1.230:7077
```

Spark Standalone Cluster Manager UI - idle

Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077
Alive Workers: 1
Cores in use: 8 Total, 0 Used
Memory in use: 30.4 GB Total, 0.0 B Used
Applications: 0 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (0 Used)	30.4 GB (0.0 B Used)

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Running spark-shell in cluster

```
spark-shell --master spark://192.168.1.230:7077 \
 --driver-memory 1g \
 --executor-memory 2g \
 --total-executor-cores 4 \
 --executor-cores 2 \
 --jars /tmp/dataset-0.9.0-SNAPSHOT-fat.jar
```

Spark Standalone Cluster Manager - 1 running application

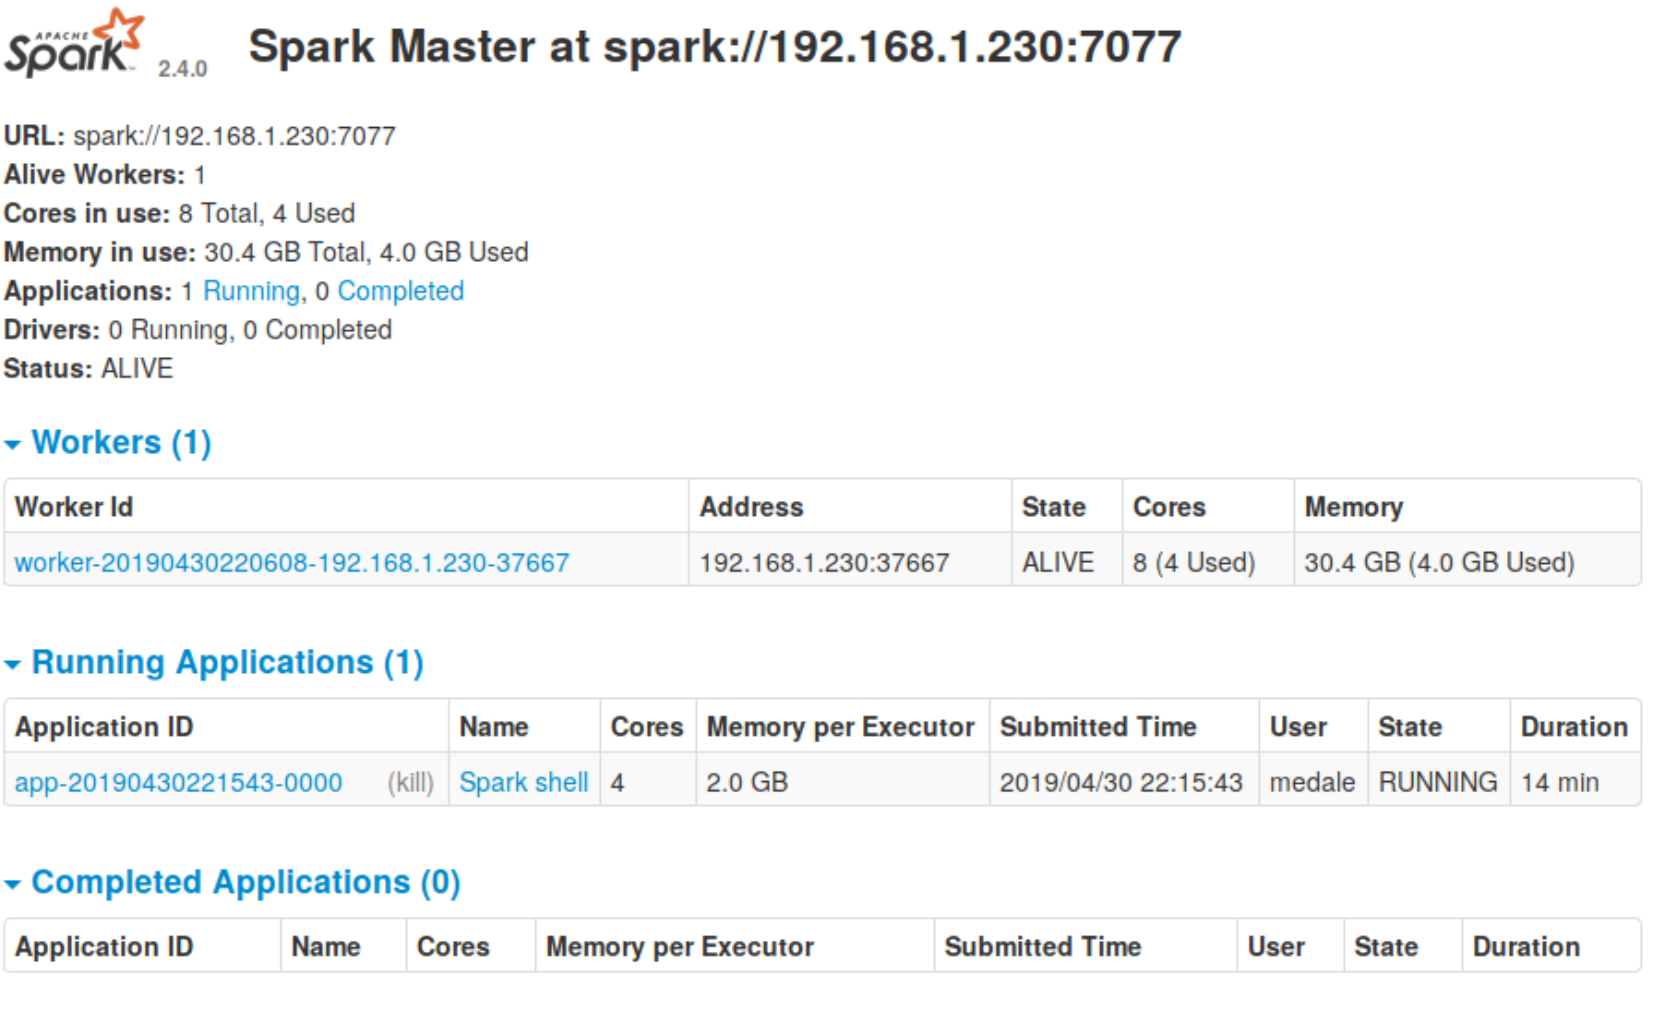

2.4.0

Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077
Alive Workers: 1
Cores in use: 8 Total, 4 Used
Memory in use: 30.4 GB Total, 4.0 GB Used
Applications: 1 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (4 Used)	30.4 GB (4.0 GB Used)

Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
app-20190430221543-0000 (kill)	Spark shell	4	2.0 GB	2019/04/30 22:15:43	medale	RUNNING	14 min

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

RDDs - Not deprecated!

```scala
object RddProcessor {

  val DefaultEventInputUrl = "file:///datasets/github/data"

  def process(sc: SparkContext, inputUrl: String): (Long, Long) = {
    val records: RDD[String] = sc.textFile(inputUrl)
    println(s"We have a total of ${records.partitions.size} partitions.")
    val total = records.count()
    val prs = records.filter(r => r.contains("PullRequestEvent"))
    val totalPrs = prs.count()
    (total, totalPrs)
  }

  def main(args: Array[String]): Unit = {
    val spark = SparkSession.builder().
      appName( name = "RddProcessor").
      getOrCreate()

    val inputUrl = if (args.size > 0) {
      args(0)
    } else {
      DefaultEventInputUrl
    }
    process(spark.sparkContext, inputUrl)
    spark.stop()
  }
}
```

And now for something completely different: Colon Cancer

- Screening saves lives!
  - Colonoscopy - talk to your doc
  - Dave Barry: A journey into my colon — and yours
- Colorectal Cancer Alliance

Questions?

- medale@asymmetrik.com
- Infrequent blog/past presentations http://uebercomputing.com/