

Data Engineering with Apache Spark

Markus Dale, medale@asymmetrik.com

May 2019

Intro, Slides And Code

- Slides: <https://github.com/medale/prez-spark-dataengineering/blob/master/presentation/SparkDataEngineering.pdf>
- Spark Data Engineering Code Examples:
<https://github.com/medale/prez-spark-dataengineering>

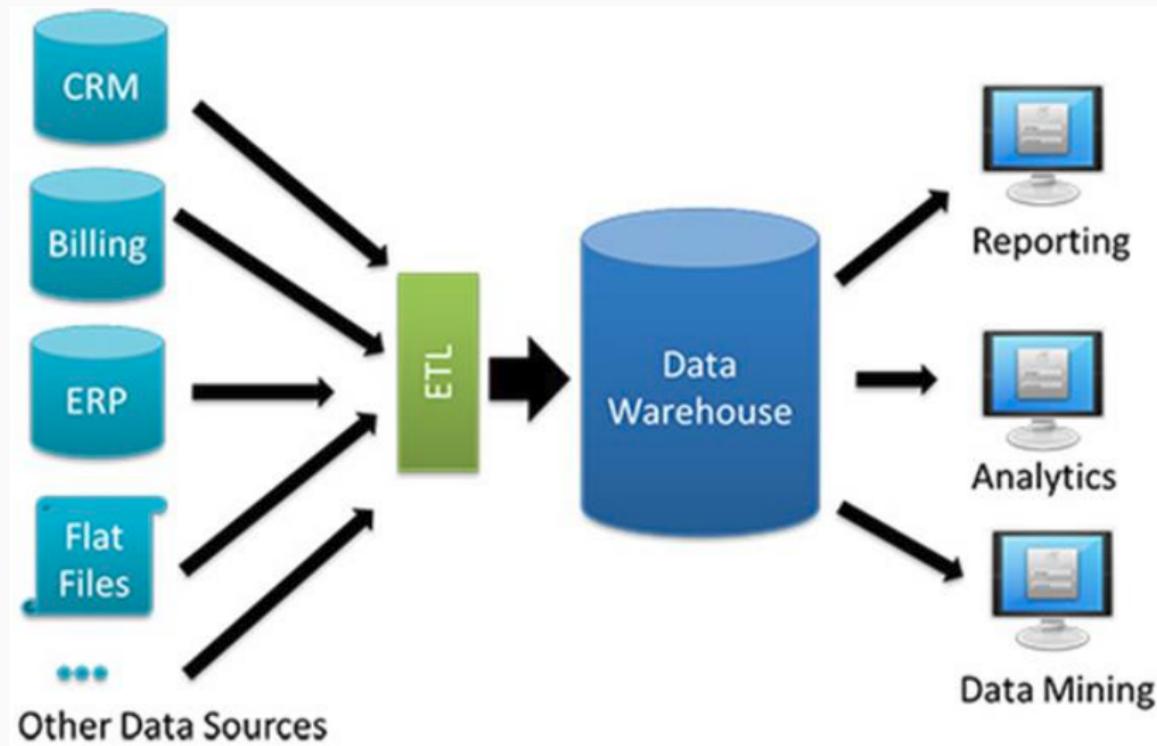
Data Science Mission

- ID malicious GitHub Pull Requests?
- Source: <https://www.gharchive.org/>

Data Engineering Mission

- <https://www.gharchive.org/>
 - 2/12/2011-12/31/2014 Timeline API (now deprecated).
 - From 1/1/2015 to now Events API.
 - Since 2015: About 20-40MB/hour * 37900 hours ~1TB
- Store: Fast, time-based access when specifying yyyy, mm, dd, hh (or prefix combination)
- Side effect: Learn more about Spark batch processing

Data engineering



Apache Spark - Big data tooling



Figure 1: Swiss Army Knife for Big Data

Apache Spark: Data engineering on small dataset



Figure 2: Laptop

Apache Spark: Data engineering for larger dataset (Vertical Scaling)




Figure 3: Beefed-up Server

Apache Spark: Data engineering for large datasets (Horizontal Scaling)

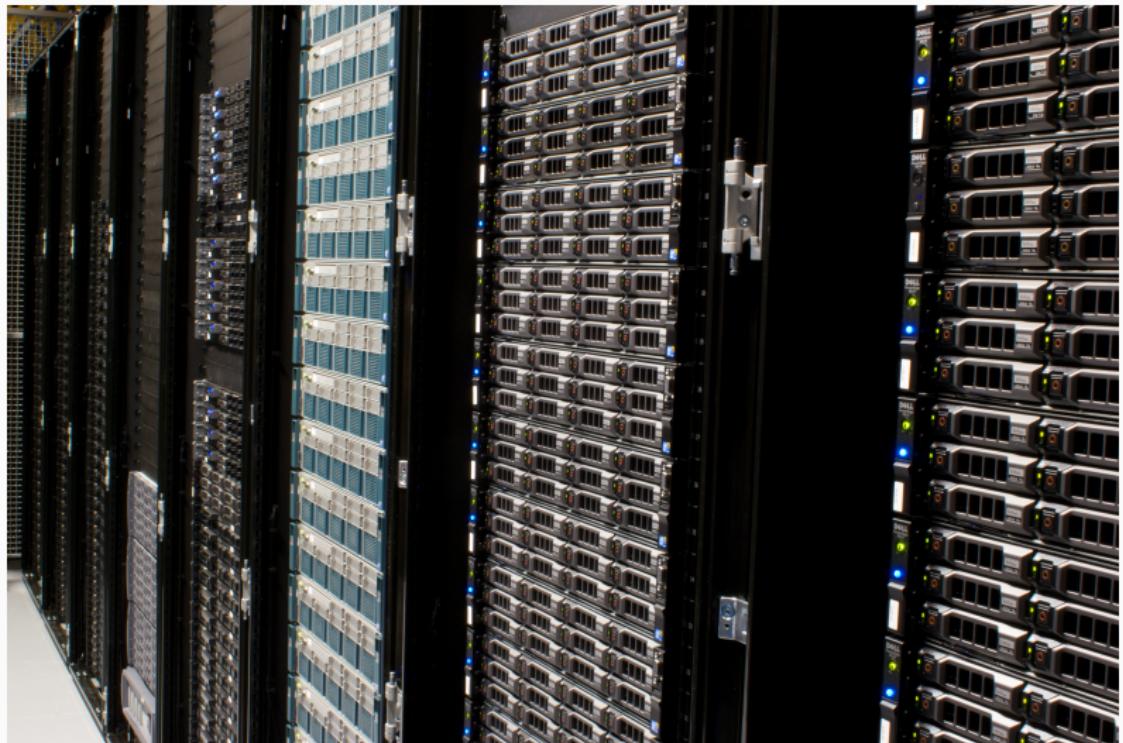


Figure 4: Multiple cooperating Servers

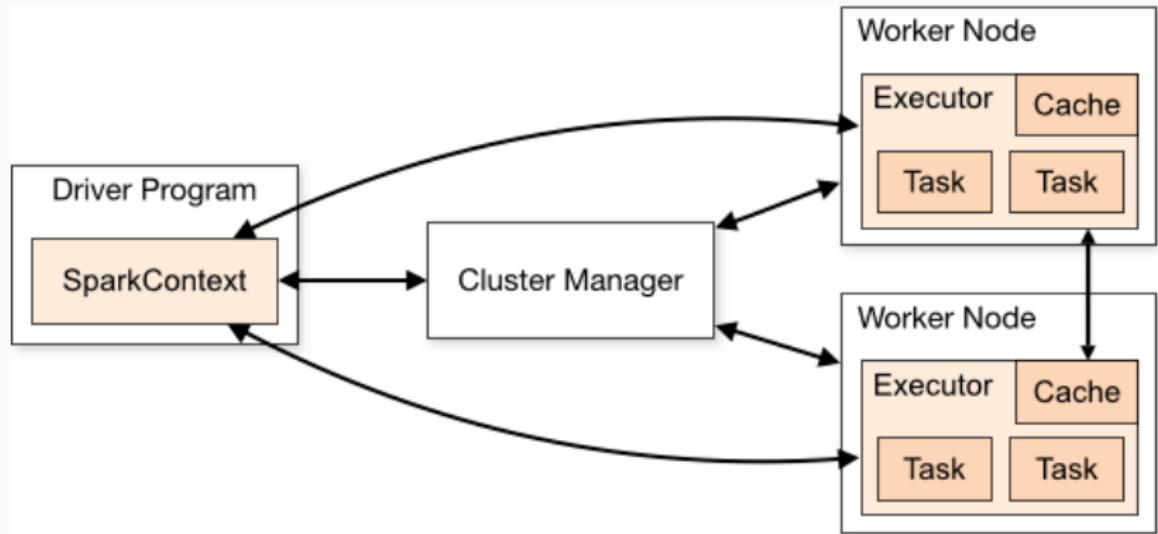
Cluster Manager - Manage cores, memory, special capabilities



Amazon EMR



Anatomy of a Spark Application



Source: Apache Spark website

Hello, Spark World!

```
import org.apache.spark.sql.SparkSession

object HelloSparkWorld {

    val RecordsUrl = "file:///datasets/github/data"

    def process(spark: SparkSession): (Long, Long) = {
        val records = spark.read.json(RecordsUrl)
        val totalEventCount = records.count()
        val prs = records.where(records("type") === "PullRequestEvent")
        val pullRequestEventCount = prs.count()
        (totalEventCount, pullRequestEventCount)
    }

    def main(args: Array[String]): Unit = {
        val spark = SparkSession.builder().appName("HelloSparkWorld").getOrCreate()
        val (total, prs) = process(spark)
        println(s"Total events: ${total}, pr events: ${prs}.")
        spark.stop()
    }
}
```

SparkSession - Gateway to the Cluster

```
def main(args: Array[String]): Unit = {  
    val spark = SparkSession.builder().  
        appName("HelloSparkWorld").  
        getOrCreate()  
    ...  
    spark.stop()  
}
```

API - SparkSession Object

SparkSession

org.apache.spark.sql

SparkSession

Related Docs: [class Sp](#)

object **SparkSession** extends [Logging](#) with [Serializable](#)

Annotations @Stable()
Source [SparkSession.scala](#)

► Linear Supertypes

Ordering [Alphabetic](#) By Inheritance

Inherited [SparkSession](#) [Serializable](#) [Serializable](#) [Logging](#) [AnyRef](#) [Any](#)
[Hide All](#) [Show All](#)

Visibility [Public](#) All

Type Members

► class **Builder** extends [Logging](#)
Builder for [SparkSession](#).

Value Members

► def **active**: [SparkSession](#)
Returns the currently active SparkSession, otherwise the default one.

► def **builder()**: [Builder](#)
Creates a [SparkSession.Builder](#) for constructing a [SparkSession](#).

API - SparkSession Class, DataFrame = Dataset[Row]

```
//spark.read - DataFrameReader  
val records: DataFrame =  
  spark.read.json(RecordsUrl)
```

spark.read: DataFrameReader - Input

- CSV
- json
- jdbc
- parquet
- text - DataFrame with “value” column
- textFile - Dataset[String]
- Third party:
 - <https://spark-packages.org>: Avro, Redshift, MongoDB...
 - Spark Cassandra Connector (DataStax github)

DataFrame Schema

```
scala> records.printSchema
root
|-- actor: struct (nullable = true)
|   |-- display_login: string (nullable = true)
|   |-- id: long (nullable = true)
...
|-- created_at: string (nullable = true)
|-- id: string (nullable = true)
|-- payload: struct (nullable = true)
|   |-- comment: struct (nullable = true)
|   |   |-- body: string (nullable = true)
...
|-- public: boolean (nullable = true)
|-- repo: struct (nullable = true)
|   |-- id: long (nullable = true)
|   |-- url: string (nullable = true)
|-- type: string (nullable = true)
```

GitHub Event Data

```
cd /datasets/github/data
wget http://data.gharchive.org/2019-04-28-0.json.gz
wget http://data.gharchive.org/2019-04-28-1.json.gz
wget http://data.gharchive.org/2019-04-28-13.json.gz
```

Preliminary Exploration

```
# creates 2019-04-28-0.json  
gunzip 2019-04-28-0.json.gz
```

```
# 95865  
wc -l 2019-04-28-0.json
```

Editor: one JSON per line

```
[{"id": "9493263421", "type": "PushEvent", "actor": {"id": 25541826, "login": "mainh", "display_login": "mainh", "gravatar_id": "", "url": "https://api.github.com/users/mainh", "avatar_url": "https://avatars.githubusercontent.com/u/25541826?"}, "repo": {"id": 150551511, "name": "Car-eye-team/JTT1078-media-server", "url": "https://api.github.com/repos/Car-eye-team/JTT1078-media-server"}, "payload": {"push_id": 3536163318, "size": 1, "distinct_size": 1, "ref": "refs/heads/master", "head": "fa8714446ae5c83c16f3125c33d51388aa6e797c", "before": "6edb3514ce440ed16d8e4cc99b8f829211f4e8a8", "commits": [{"sha": "fa8714446ae5c83c16f3125c33d51388aa6e797c", "author": {"name": "Main", "email": "a3a6910872d9e29b8a7d713ee6dd76d73f01233c@qq.com"}, "message": "Add files via upload\n\n更新架构图", "distinct": true, "url": "https://api.github.com/repos/Car-eye-team/JT1078-media-server/commits/fa8714446ae5c83c16f3125c33d51388aa6e797c"}]}, "public": true, "created_at": "2019-04-23T12:00:00Z", "org": {"id": 36096525, "login": "Car-eye-team", "gravatar_id": "", "url": "https://api.github.com/orgs/Car-eye-team", "avatar_url": "https://avatars.githubusercontent.com/u/36096525?"}}]  
{"id": "9493263424", "type": "WatchEvent", "actor": {"id": 6560033, "login": "voby", "display_login": "voby", "gravatar_id": "", "url": "https://api.github.com/users/voby", "avatar_url": "https://avatars.githubusercontent.com/u/6560033?"}, "repo": {"id": 74293321, "name": "sveltejs/svelte", "url": "https://api.github.com/repos/sveltejs/svelte"}, "payload": {"action": "started"}, "public": true, "created_at": "2019-04-23T12:00:00Z", "org": {"id": 23617963, "login": "sveltejs", "gravatar_id": "", "url": "https://api.github.com/orgs/sveltejs", "avatar_url": "https://avatars.githubusercontent.com/u/23617963?"}}}  
{"id": "9493263426", "type": "PullRequestEvent", "actor": {"id": 300805, "login": "D4nte", "display_login": "D4nte", "gravatar_id": "", "url": "https://api.github.com/users/D4nte", "avatar_url": "https://avatars.githubusercontent.com/u/300805?"}, "repo": {"id": 145489725, "name": "coblox/bitcoinrpc-rust-client", "url": "https://api.github.com/repos/coblox/bitcoinrpc-rust-client"}, "payload": {"action": "opened", "number": 54, "pull_request": {"url": "https://api.github.com/repos/coblox/bitcoinrpc-rust-client/pull/54?"}}}
```

Pretty Print One Record?

```
# default 1000 lines - xaa, xab...
split 2019-04-28-0.json
mkdir temp
cd temp

# 1 file per line
split -1 .../xaa

# xac's a PullRequestEvent
python -m json.tool < xac > pretty.json
```

Open pretty.json in Atom - PullRequestEvent

```
        "organizations_url": "https://api.github.com/users/D4nte/orgs",
        "received_events_url": "https://api.github.com/users/D4nte/received_events",
        "repos_url": "https://api.github.com/users/D4nte/repos",
        "site_admin": false,
        "starred_url": "https://api.github.com/users/D4nte/starred{/owner}",
        "subscriptions_url": "https://api.github.com/users/D4nte/subscriptions",
        "type": "User",
        "url": "https://api.github.com/users/D4nte"
    }
},
{
    "public": true,
    "repo": {
        "id": 145489725,
        "name": "cloblox/bitcoinrpc-rust-client",
        "url": "https://api.github.com/repos/cloblox/bitcoinrpc-rust-client"
    },
    "type": "PullRequestEvent"
}
```

Starting Spark Standalone Cluster Manager

```
# Start on master  
$SPARK_HOME/sbin/start-master.sh --host 192.168.1.232  
  
# Start one or more workers  
$SPARK_HOME/sbin/start-slave.sh spark://192.168.1.232:7077
```

Spark Standalone Cluster Manager UI - idle



Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077

Alive Workers: 1

Cores in use: 8 Total, 0 Used

Memory in use: 30.4 GB Total, 0.0 B Used

Applications: 0 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (0 Used)	30.4 GB (0.0 B Used)

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Completed Applications (0)


Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Running spark-shell in cluster

```
spark-shell --master spark://192.168.1.232:7077 \
--driver-memory 1g \
--executor-memory 2g \
--total-executor-cores 4 \
--executor-cores 2 \
--jars /tmp/dataset-0.9.0-SNAPSHOT-fat.jar
```

Spark Shell Startup

```
ng bulletin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For
Spark context Web UI available at http://markus-desktop-u.
Spark context available as 'sc' (master = spark://192.168.0
Spark session available as 'spark'.
Welcome to
```



```
Using Scala version 2.11.12 (Java HotSpot(TM) 64-Bit Server
Type in expressions to have them evaluated.
Type :help for more information.
```

```
scala> :quit
```

Spark Standalone Cluster Manager - 1 running application



Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077

Alive Workers: 1

Cores in use: 8 Total, 4 Used

Memory in use: 30.4 GB Total, 4.0 GB Used

Applications: 1 [Running](#), 0 [Completed](#)

Drivers: 0 Running, 0 Completed

Status: ALIVE

▼ Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (4 Used)	30.4 GB (4.0 GB Used)

▼ Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
app-20190430221543-0000	(kill) Spark shell	4	2.0 GB	2019/04/30 22:15:43	medale	RUNNING	14 min

▼ Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

spark-shell auto-imports

```
scala> :imports
1) import org.apache.spark.SparkContext._
2) import spark.implicits._
3) import spark.sql
4) import org.apache.spark.sql.functions._
```

Data Exploration - schema and counting

```
val RecordsUrl = "file:///datasets/github/data"
val records = spark.read.json(RecordsUrl)
records.printSchema
//...
// |-- repo: struct (nullable = true)
// |   |-- id: long (nullable = true)
// |   |-- name: string (nullable = true)
// |   |-- url: string (nullable = true)
// |   |-- type: string (nullable = true)

records.count()
//147,374
```

Spark Application UI - Jobs, stages, tasks



Spark shell application UI

Jobs Stages Storage Environment Executors SQL

Spark Jobs (?)

User: medale
Total Uptime: 2.5 min
Scheduling Mode: FIFO
Completed Jobs: 2

▶ Event Timeline

▼ Completed Jobs (2)

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at <console>:26 count at <console>:26	2019/05/02 10:01:16	0.6 s	2/2	4/4
0	json at <console>:25 json at <console>:25	2019/05/02 10:00:58	4 s	1/1	3/3

Spark Application UI - Stages



Spark shell application UI

Jobs Stages Storage Environment Executors SQL

Stages for All Jobs

Completed Stages: 3

▼ Completed Stages (3)

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	count at <console>:26 +details	2019/05/02 10:01:16	76 ms	1/1			177.0 B	
1	count at <console>:26 +details	2019/05/02 10:01:16	0.5 s	3/3	46.6 MB			177.0 B
0	json at <console>:25 +details	2019/05/02 10:00:58	4 s	3/3	46.6 MB			

Spark Application UI - Stage details

Jobs Stages Storage Environment Executors SQL

Details for Stage 1 (Attempt 0)


Total Time Across All Tasks: 1 s

Locality Level Summary: Process local: 3

Input Size / Records: 46.6 MB / 147374

Shuffle Write: 177.0 B / 3

▼ DAG Visualization



► Show Additional Metrics


► Event Timeline

Summary Metrics for 3 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
--------	-----	-----------------	--------	-----------------	-----

Job 1 - Stages 1 and 2 DAG

DAG Visualization



One Job = n lazy transformations, 1 action

```
//job0 (read twice - once for schema, once for count)
val records = spark.read.json(RecordsUrl)
//lazy transformation
records.cache()
//action 1 - end of job 1
records.count()
records.printSchema

//2 lazy transformations - immutable datasets
val types = records.select("type")
val distinctTypes = types.distinct()

//one eager action - job 2
distinctTypes.show()
+-----+
|      type |
+-----+
|      PushEvent |
...
|      PullRequestEvent |
+-----+
```

Spark Application UI - Storage (caching)

APACHE  2.4.0

Jobs Stages Storage Environment Executors SQL

Spark shell application UI

Storage

▼ RDDs

ID	RDD Name	Storage Level	Cached Partitions	Fraction Cached	Size in Memory	Size on Disk
11	FileScan json [actor#6,created_at#7,id#8,org#9,payload#10,public#11,repo#12,type#13] Batched: false, Format: JSON, Location: InMemoryFileIndex[file:/datasets/github/data], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<actor:struct<avatar_url:string,display_login:string,gravatar_id:string,id:bigint,login:str...]	Memory Deserialized 1x Replicated	3	100%	312.1 MB	0.0 B

Default file system/file system URLs

```
val hc = spark.sparkContext.hadoopConfiguration  
hc.get("fs.defaultFS")  
//res4: String = file:///
```

```
spark.read.json("/datasets/github/data")  
spark.read.json("file:///datasets/github/data")
```

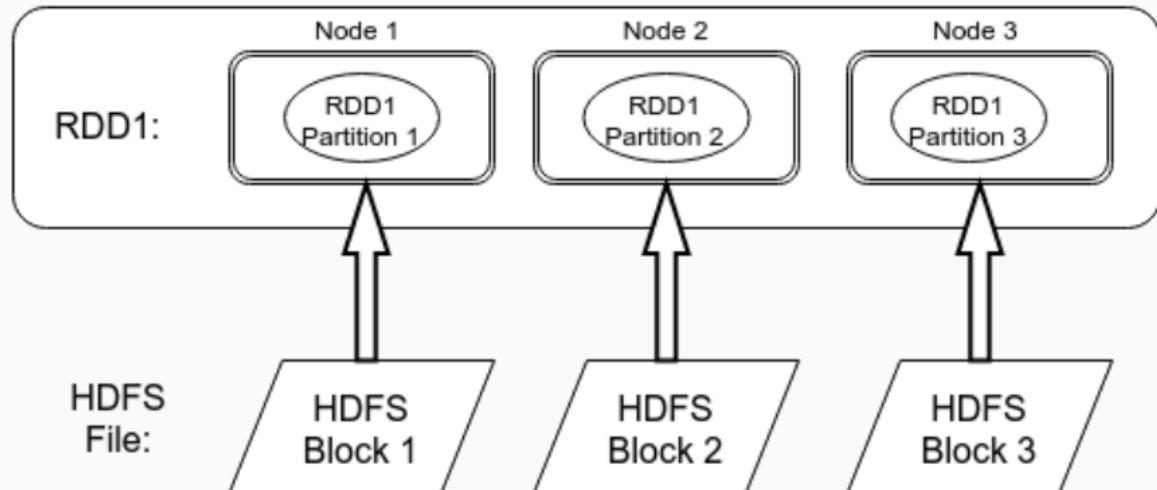
```
//hdfs://<name-node>:port  
//s3a://<bucket-name>
```

Input partitions - splittable file?

- Splittable: bzip2, parquet, avro
- Non-splittable: gzip (1 task per file)

```
hc.get("dfs.block.size")  
//res6: String = 134217728
```

Resilient Distributed Datasets (RDDs)



RDDs - Not deprecated!

```
object RddProcessor {  
  
    val DefaultEventInputUrl = "file:///datasets/github/data"  
  
    def process(sc: SparkContext, inputUrl: String): (Long, Long) = {  
        val records: RDD[String] = sc.textFile(inputUrl)  
        println(s"We have a total of ${records.partitions.size} partitions.")  
        val total = records.count()  
        val prs = records.filter(r => r.contains("PullRequestEvent"))  
        val totalPrs = prs.count()  
        (total, totalPrs)  
    }  
  
    def main(args: Array[String]): Unit = {  
        val spark = SparkSession.builder().  
            appName(name = "RddProcessor").  
            getOrCreate()  
  
        val inputUrl = if (args.size > 0) {  
            args(0)  
        } else {  
            DefaultEventInputUrl  
        }  
        process(spark.sparkContext, inputUrl)  
        spark.stop()  
    }  
}
```

Datasets/DataFrames compiled to RDDs

- Catalyst query optimizer for built-in functions
- Project Tungsten - memory management
 - Row storage (Apache Arrow)
 - Encoders for Dataset objects (`spark.implicits._`)

Data Exploration - event type distribution

```
val prs = records.where(records("type") === "PullRequestEvent")
val pullRequestEventCount = prs.count()
//6699

//import spark.implicits._ - auto-imported by shell
val typeCounts = records.groupBy($"type").
    count.orderBy($"count".desc)
typeCounts.show
+-----+-----+
|      type | count |
+-----+-----+
| PushEvent | 82518 |
| CreateEvent | 18313 |
| WatchEvent | 16868 |
| PullRequestEvent | 6699 |
...
...
```

Narrow vs. wide transformations

- narrow: all data from one partition
- wide: data from multiple parent partitions (shuffle)

Shuffle Partitions - Default

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
4	show at <console>:26 show at <console>:26	2019/05/02 10:54:53	0.8 s	2/2	203/203

Setting Shuffle Partitions

```
spark.conf.set("spark.sql.shuffle.partitions", "10")
```

Shuffle Partitions Optimized

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	show at <console>:26 show at <console>:26	2019/05/02 11:04:14	0.1 s	2/2	13/13

API - Some Dataset Transformations

- select
- where
- distinct
- limit
- orderBy
- join

API - Some Dataset Actions

- collect
- count
- take(n)
- head
- write - DataFrameWriter

DataFrameWriter

- CSV
- JDBC
- Parquet
- Text

Just the PullRequestEvents and their schema

```
val texts = spark.read.text(RecordsUrl)
val prsText =
  texts.where($"value".contains("PullRequestEvent"))

val reparteds = prsText.repartition(2)
reparteds.write.text("file:///datasets/github/prs")
```

Writing partitions - Output directory (57MB/file)

```
ls
```

```
_SUCCESS
```

```
part-00000-9f85464d-c4e8-4b55-9d95-580acb3f30cd-c000.txt
```

```
part-00001-9f85464d-c4e8-4b55-9d95-580acb3f30cd-c000.txt
```

API - Column

- +, -, *, %
- ==, !=, >, <, ...
- asc, desc
- startsWith, contains, endsWith, like, rlike
- isNull, isNaN, isIn

API - functions

- array functions, explode
- date/time functions
- math, string

Date Exploration

```
val prs = spark.read.json("file:///datasets/github/prs")
//prs: org.apache.spark.sql.DataFrame = [actor: struct<avatar_url: string,
//display_login: string ... 4 more fields>, created_at: string

prs.select("created_at").take(10)
//res15: Array[org.apache.spark.sql.Row] =
//Array([2019-04-28T01:37:47Z],
//[2019-04-28T01:30:24Z], [2019-04-28T01:23:30Z],...
```

Adding year, month, day, hour columns

```
val ymdhPrs = prs.withColumn("year", year($"created_at")).  
  withColumn("month", month($"created_at")).  
  withColumn("day", dayofmonth($"created_at")).  
  withColumn("hour", hour($"created_at"))  
  
ymdhPrs.printSchema  
...  
|-- type: string (nullable = true)  
|-- year: integer (nullable = true)  
|-- month: integer (nullable = true)  
|-- day: integer (nullable = true)  
|-- hour: integer (nullable = true)
```

Saving to Parquet with partitioning columns

```
ymdhPrs.write.partitionBy("year","month","day","hour").  
    parquet("file:///datasets/github/prs-ymdh")
```

```
/datasets/github/prs-ymdh/year=2019/month=4/day=28/hour=9  
part-00000-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet  
part-00001-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet  
part-00002-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet
```

```
/datasets/github/prs-ymdh/year=2019/month=4/day=28/hour=20  
part-00001-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet  
part-00002-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet  
part-00003-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet
```

```
/datasets/github/prs-ymdh/year=2019/month=4/day=28/hour=21  
part-00000-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet  
part-00003-afd2a04d-bdeb-48fe-9e43-12dc4fb3535a.c000.snappy.parquet
```

Reading from Parquet - schema and predicate pushdown

```
val mprs = spark.read.parquet("file:///datasets/github/prs-ymdh")
val oneHour = mprs.where("year = 2019 AND month = 04 AND " +
    "day = 28 AND hour = 21")

oneHour.explain(true)
...
== Optimized Logical Plan ==
Filter (((((isnotnull(day#493) && isnotnull(year#491)) &&
    isnotnull(month#492))
&& isnotnull(hour#494)) && (year#491 = 2019)) && (month#492 = 4)) && (day#493
    = 28))
&& (hour#494 = 21))
+- Rela-
    tion[actor#483,created_at#484,id#485,org#486,payload#487,public#488,repo#489,
    type#490,year#491,month#492,day#493,hour#494] parquet

oneHour.write.parquet("/datasets/github/one")
mprs.write.parquet("/datasets/github/all")
```

One Hour Completed Stage

▼ Completed Stages (1)

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
30	parquet at <console>:26 +details	2019/05/02 12:56:50	2 s	2/2	7.3 MB	6.9 MB		

All Completed Stage

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
29	count at <console>:26 +details	2019/05/02 12:51:14	11 ms	1/1			236.0 B	
28	count at <console>:26 +details	2019/05/02 12:51:14	71 ms	4/4	1935.6 KB			236.0 B

Storage size compression

- original gz - 47MB
- text only - 115MB
- parquet - 27.3MB

Apache Spark?



Figure 5: Questions?

And now for something completely different: Colon Cancer



- Screening saves lives!
 - Colonoscopy - talk to your doc
 - Dave Barry: A journey into my colon — and yours
- Colorectal Cancer Alliance

Questions?



- medale@asymmetrik.com
- <https://github.com/medale/prez-spark-dataengineering/blob/master/presentation/SparkDataEngineering.pdf>