

Data Engineering with Apache Spark

Markus Dale, medale@asymmetrik.com

May 2019

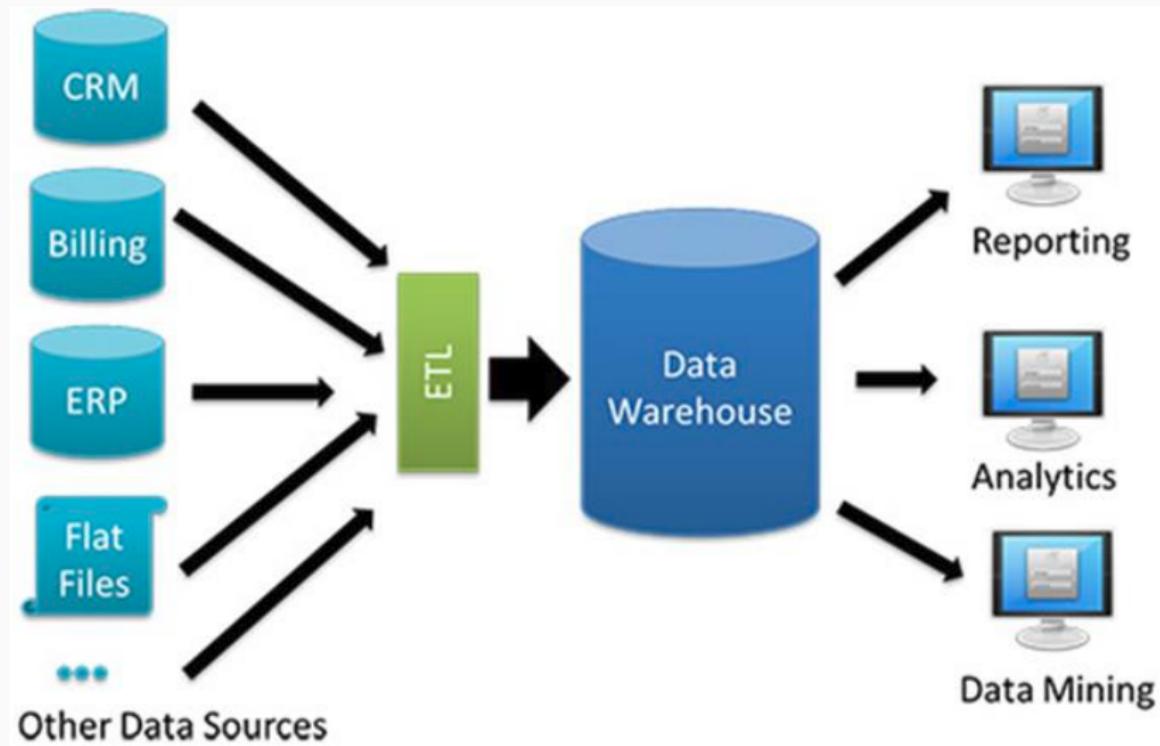
Intro, Slides And Code

- Slides: <https://github.com/medale/prez-spark-dataengineering/blob/master/presentation/SparkDataEngineering.pdf>
- Scala Spark Code Examples:
<https://github.com/medale/prez-spark-dataengineering>

Goals

- Intro to Spark Scala Dataset API for data engineering
 - At scale data exploration
 - At scale ETL (Extract Transform Load)
 - Storage formats
 - Spark abstractions

Data engineering



Apache Spark: Data engineering on small dataset



Figure 1: Laptop

Apache Spark: Data engineering for larger dataset (Vertical Scaling)

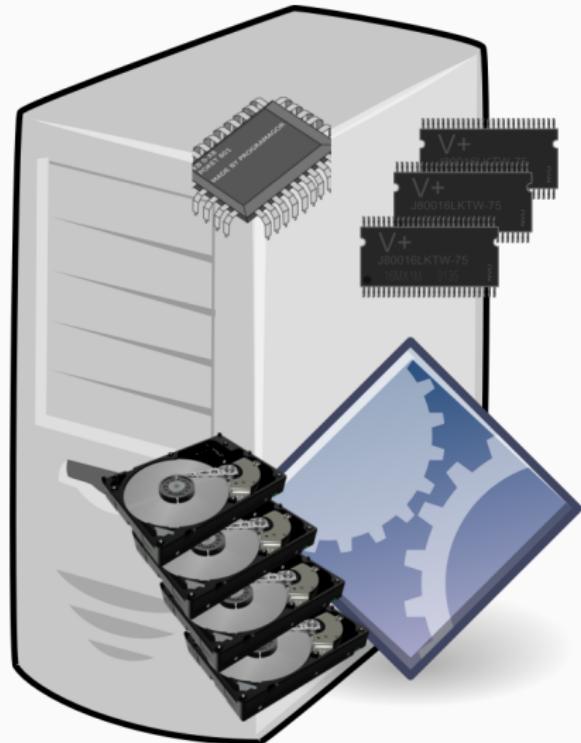


Figure 2: Beefed-up Server

Data engineering for large datasets (Horizontal Scaling)

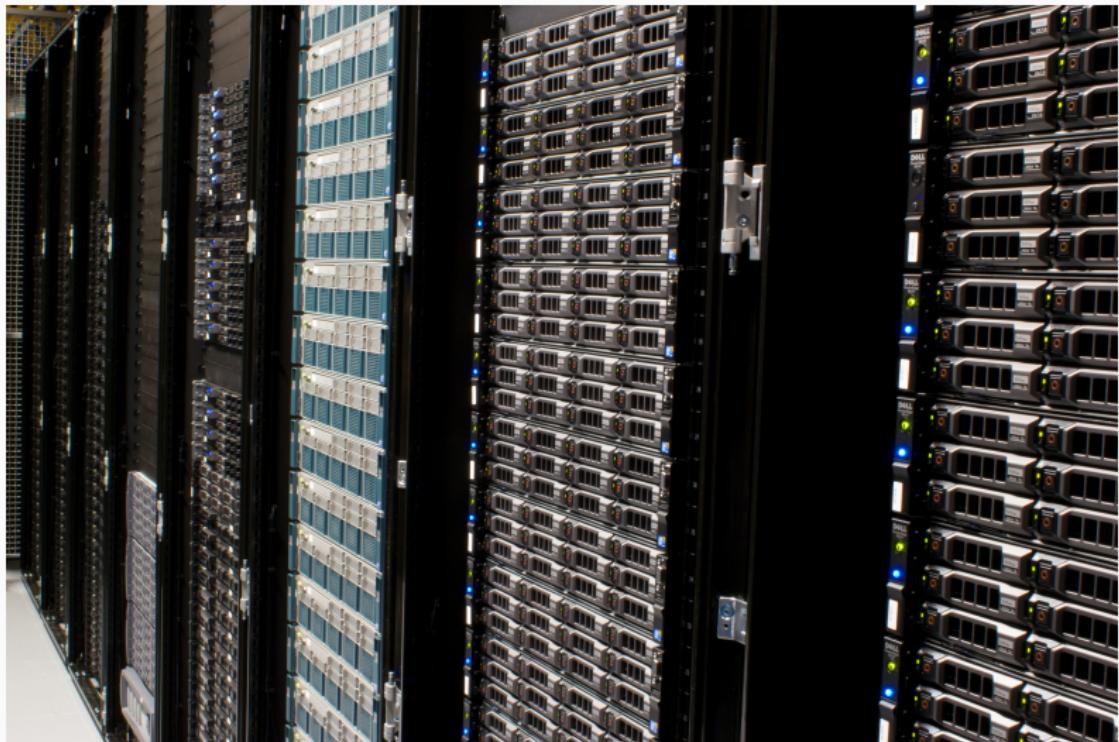


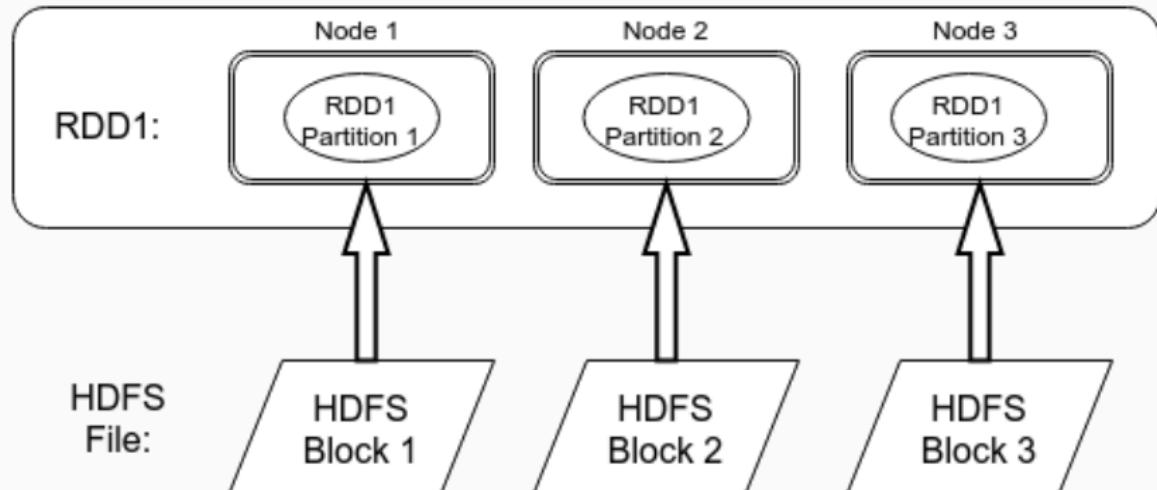
Figure 3: Multiple cooperating Servers



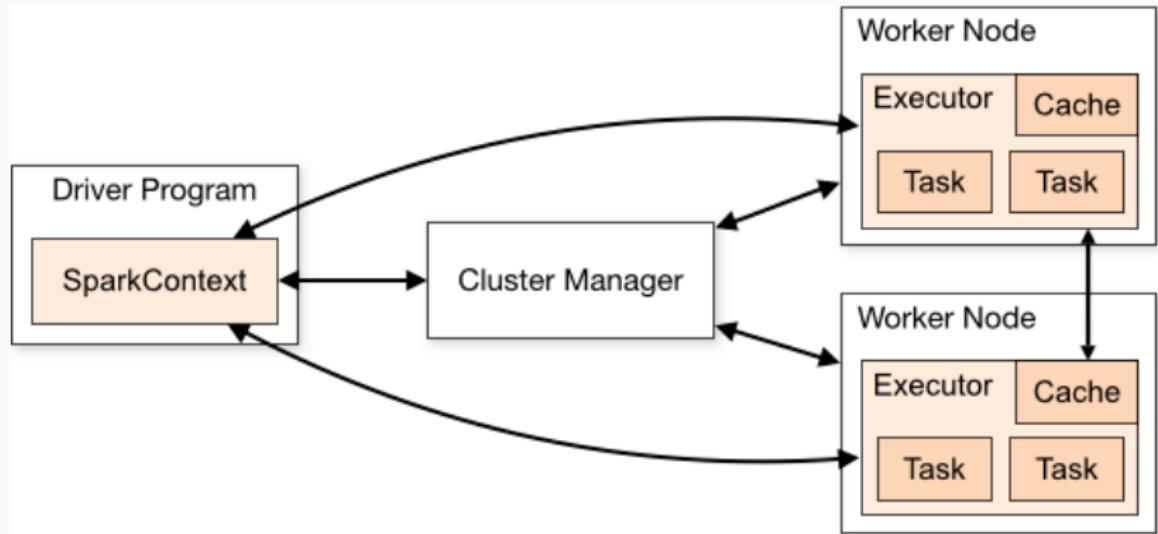
Amazon EMR



Resilient Distributed Datasets (RDDs)



Anatomy of a Spark Application



Source: Apache Spark website

Hello, Spark World!

```
import org.apache.spark.sql.SparkSession

object HelloSparkWorld {

    def process(spark: SparkSession): (Long, Long) = {
        val records = spark.read.json(path = "file:///datasets/github/data")
        records.cache()
        val totalEventCount = records.count()

        val prs = records.where(records("type") === "PullRequestEvent")
        val pullRequestEventCount = prs.count()

        records.unpersist()
        (totalEventCount, pullRequestEventCount)
    }

    def main(args: Array[String]): Unit = {
        val spark = SparkSession.builder().
            appName(name = "HelloSparkWorld").
            getOrCreate()
        process(spark)
    }
}
```

Starting Spark Standalone Cluster Manager

```
# Start on master  
$SPARK_HOME/sbin/start-master.sh --host 192.168.1.232  
  
# Start one or more workers  
$SPARK_HOME/sbin/start-slave.sh spark://192.168.1.232:7077
```

Spark Standalone Cluster Manager UI - idle



Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077

Alive Workers: 1

Cores in use: 8 Total, 0 Used

Memory in use: 30.4 GB Total, 0.0 B Used

Applications: 0 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (0 Used)	30.4 GB (0.0 B Used)

Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

Running spark-shell in cluster

```
spark-shell --master spark://192.168.1.232:7077 \
--driver-memory 1g \
--executor-memory 2g \
--total-executor-cores 4 \
--executor-cores 2 \
--jars /tmp/dataset-0.9.0-SNAPSHOT-fat.jar
```

Spark Shell Startup

```
ng bulletin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For
Spark context Web UI available at http://markus-desktop-u.
Spark context available as 'sc' (master = spark://192.168.0
Spark session available as 'spark'.
Welcome to
```



```
Using Scala version 2.11.12 (Java HotSpot(TM) 64-Bit Server
Type in expressions to have them evaluated.
Type :help for more information.
```

```
scala> :quit
```

Spark Standalone Cluster Manager - 1 running application



Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077

Alive Workers: 1

Cores in use: 8 Total, 4 Used

Memory in use: 30.4 GB Total, 4.0 GB Used

Applications: 1 [Running](#), 0 [Completed](#)

Drivers: 0 Running, 0 Completed

Status: ALIVE

▼ Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (4 Used)	30.4 GB (4.0 GB Used)

▼ Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
app-20190430221543-0000	(kill) Spark shell	4	2.0 GB	2019/04/30 22:15:43	medale	RUNNING	14 min

▼ Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration

HelloSparkWorld in spark-shell

```
scala> import com.uebercomputing.HelloSparkWorld  
import com.uebercomputing.HelloSparkWorld  
  
scala> HelloSparkWorld.process(spark)  
res0: (Long, Long) = (147374,6699)
```

Spark Application UI - Jobs, stages, tasks



Spark shell application UI

Jobs Stages Storage Environment Executors SQL

Spark Jobs (?)

User: medale
Total Uptime: 7.5 min
Scheduling Mode: FIFO
Completed Jobs: 3

▶ Event Timeline

▼ Completed Jobs (3)

Job Id ▾	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	count at HelloSparkWorld.scala:16 count at HelloSparkWorld.scala:16	2019/04/30 22:54:02	0.3 s	2/2	4/4
1	count at HelloSparkWorld.scala:13 count at HelloSparkWorld.scala:13	2019/04/30 22:53:57	4 s	2/2	4/4
0	json at HelloSparkWorld.scala:11 json at HelloSparkWorld.scala:11	2019/04/30 22:53:53	3 s	1/1	3/3

Job - n lazy transformations, 1 action

```
//job 0 - list files, infer schema
val records = spark.read.json("file:///datasets/github/data")
//transformation
records.cache()
//action - job 1
val totalEventCount = records.count()

//transformation - datasets are immutable!
val prs = records.where(records("type") === "PullRequestEvent")
//action - job 2
val pullRequestEventCount = prs.count()
```

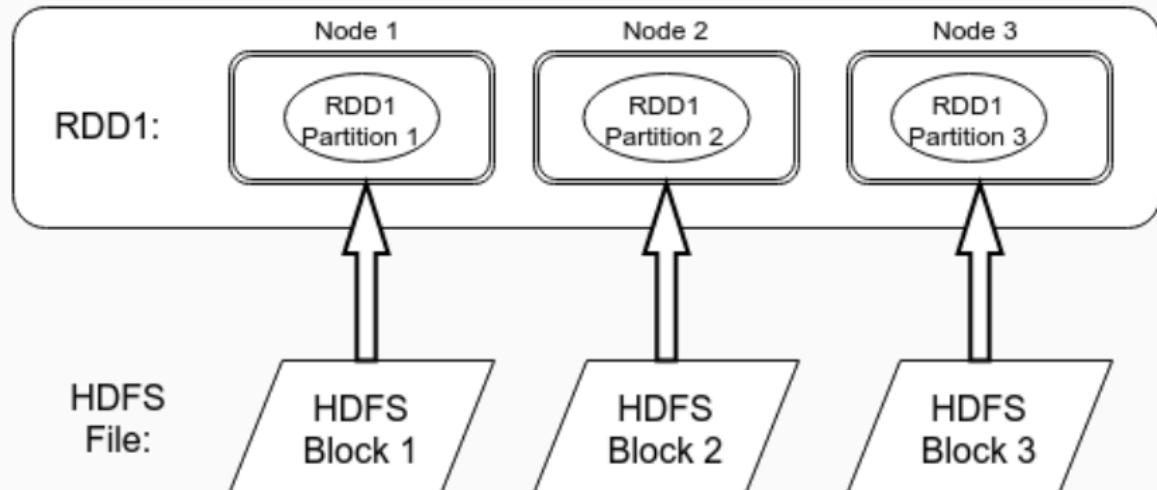
Job 0 - Stages, tasks, partitions

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at HelloSparkWorld.scala:13 count at HelloSparkWorld.scala:13	2019/04/30 22:53:57	4 s	2/2	4/4
0	json at HelloSparkWorld.scala:11 json at HelloSparkWorld.scala:11	2019/04/30 22:53:53	3 s	1/1	3/3

Job 0 - Stage 0, tasks, partitions

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0	Json at HelloSparkWorld.scala:11 +details	2019/04/30 22:53:53	3 s	3/3	46.6 MB			

Input partitions - splittable file?



Job 1 - count

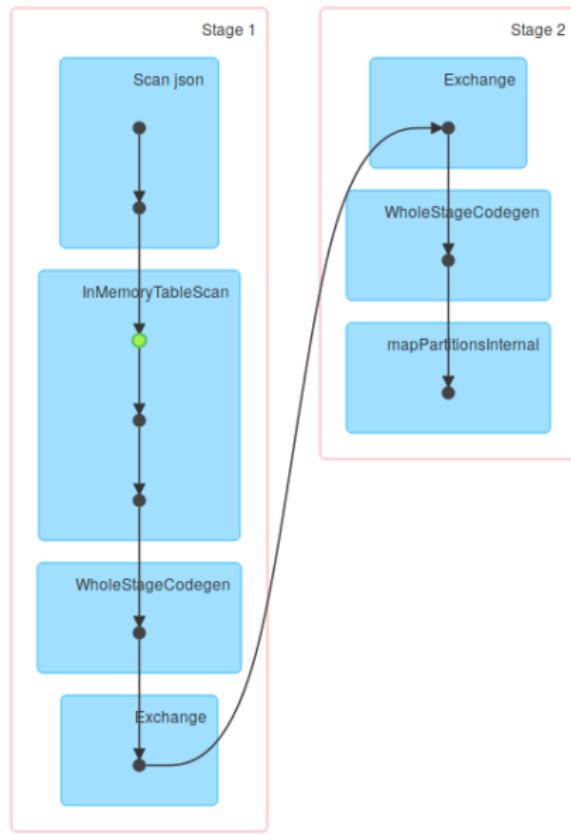
Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at HelloSparkWorld.scala:13 count at HelloSparkWorld.scala:13	2019/04/30 22:53:57	4 s	2/2	4/4
0	json at HelloSparkWorld.scala:11 json at HelloSparkWorld.scala:11	2019/04/30 22:53:53	3 s	1/1	3/3

Job 1 - count Stages 1 and 2

Stage Id ▾	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	count at HelloSparkWorld.scala:13 +details	2019/04/30 22:54:02	69 ms	1/1			177.0 B	
1	count at HelloSparkWorld.scala:13 +details	2019/04/30 22:53:57	4 s	3/3	46.6 MB			177.0 B

Job 1 - Stages 1 and 2 DAG

DAG Visualization



RDDs - Not deprecated!

```
object RddProcessor {  
  
    val DefaultEventInputUrl = "file:///datasets/github/data"  
  
    def process(sc: SparkContext, inputUrl: String): (Long, Long) = {  
        val records: RDD[String] = sc.textFile(inputUrl)  
        println(s"We have a total of ${records.partitions.size} partitions.")  
        val total = records.count()  
        val prs = records.filter(r => r.contains("PullRequestEvent"))  
        val totalPrs = prs.count()  
        (total, totalPrs)  
    }  
  
    def main(args: Array[String]): Unit = {  
        val spark = SparkSession.builder().  
            appName(name = "RddProcessor").  
            getOrCreate()  
  
        val inputUrl = if (args.size > 0) {  
            args(0)  
        } else {  
            DefaultEventInputUrl  
        }  
        process(spark.sparkContext, inputUrl)  
        spark.stop()  
    }  
}
```

And now for something completely different: Colon Cancer



- Screening saves lives!
 - Colonoscopy - talk to your doc
 - Dave Barry: A journey into my colon — and yours
- Colorectal Cancer Alliance

Questions?



- medale@asymmetrik.com
- <https://github.com/medale/prez-spark-dataengineering/blob/master/presentation/SparkDataEngineering.pdf>