

Data Engineering with Apache Spark

Markus Dale, medale@asymmetrik.com

May 2019

- Slides: <https://github.com/medale/>
- Scala Spark Code Examples: <https://github.com/medale/>






Figure 1: Laptop

Data engineering for larger dataset (Vertical Scaling)




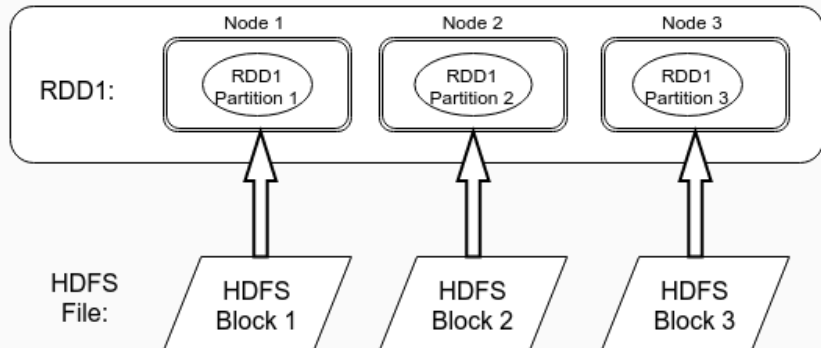
Figure 2: Beefed-up Server

Data engineering for large datasets (Horizontal Scaling)






Resilient Distributed Datasets (RDDs)



Anatomy of a Spark Application



Source: Apache Spark website

Hello, Spark World!

```
import org.apache.spark.sql.SparkSession

object HelloWorld {

  def process(spark: SparkSession): (Long,Long) = {
    val records = spark.read.json( path = "file:///datasets/github/data")
    records.cache()
    val totalEventCount = records.count()

    val prs = records.where(records("type") === "PullRequestEvent")
    val pullRequestEventCount = prs.count()

    records.unpersist()
    (totalEventCount, pullRequestEventCount)
  }

  def main(args: Array[String]): Unit = {
    val spark = SparkSession.builder().
      appName( name = "HelloSparkWorld").
      getOrCreate()
    process(spark)
  }
}
```

Starting Spark Standalone Cluster Manager

```
# Start on master
```

```
$SPARK_HOME/sbin/start-master.sh --host 192.168.1.230
```

```
# Start one or more workers
```

```
$SPARK_HOME/sbin/start-slave.sh spark://192.168.1.230:7077
```

Spark Standalone Cluster Manager UI - idle



Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077

Alive Workers: 1

Cores in use: 8 Total, 0 Used

Memory in use: 30.4 GB Total, 0.0 B Used

Applications: 0 [Running](#), 0 [Completed](#)

Drivers: 0 Running, 0 Completed

Status: ALIVE

▼ Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (0 Used)	30.4 GB (0.0 B Used)

▼ Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

▼ Completed Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

```
spark-shell --master spark://192.168.1.230:7077 \  
--driver-memory 1g \  
--executor-memory 2g \  
--total-executor-cores 4 \  
--executor-cores 2 \  
--jars /tmp/dataset-0.9.0-SNAPSHOT-fat.jar
```

Spark Shell Startup

```

ing builtin-java classes where applicable
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For
Spark context Web UI available at http://markus-desktop-u.
Spark context available as 'sc' (master = spark://192.168.
Spark session available as 'spark'.
Welcome to

```

[illegible]

```
Using Scala version 2.11.12 (Java HotSpot(TM) 64-Bit Server VM)
Type in expressions to have them evaluated.
Type :help for more information.
```

```
scala> :quit
```

Spark Standalone Cluster Manager - 1 running application



Spark Master at spark://192.168.1.230:7077

URL: spark://192.168.1.230:7077

Alive Workers: 1

Cores in use: 8 Total, 4 Used

Memory in use: 30.4 GB Total, 4.0 GB Used

Applications: 1 [Running](#), 0 [Completed](#)

Drivers: 0 Running, 0 Completed

Status: ALIVE

Workers (1)

Worker Id	Address	State	Cores	Memory
worker-20190430220608-192.168.1.230-37667	192.168.1.230:37667	ALIVE	8 (4 Used)	30.4 GB (4.0 GB Used)

Running Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
app-20190430221543-0000 (kill)	Spark shell	4	2.0 GB	2019/04/30 22:15:43	medale	RUNNING	14 min

Completed Applications (0)


Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

HelloSparkWorld in spark-shell

```
scala> import com.uebercomputing.HelloSparkWorld  
import com.uebercomputing.HelloSparkWorld
```

```
scala> HelloSparkWorld.process(spark)  
res0: (Long, Long) = (147374,6699)
```


Spark Application UI - Jobs, stages, tasks

 2.4.0

Spark shell application UI

Jobs | Stages | Storage | Environment | Executors | SQL

Spark Jobs (?)

User: medale
Total Uptime: 7.5 min
Scheduling Mode: FIFO
Completed Jobs: 3

► Event Timeline

▼ Completed Jobs (3)

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
2	count at HelloSparkWorld.scala:16 count at HelloSparkWorld.scala:16	2019/04/30 22:54:02	0.3 s	2/2	4/4
1	count at HelloSparkWorld.scala:13 count at HelloSparkWorld.scala:13	2019/04/30 22:53:57	4 s	2/2	4/4
0	json at HelloSparkWorld.scala:11 json at HelloSparkWorld.scala:11	2019/04/30 22:53:53	3 s	1/1	3/3

Job - n lazy transformations, 1 action

```
//job 0 - list files, infer schema
val records = spark.read.json("file:///datasets/github/data")
//transformation
records.cache()
//action - job 1
val totalEventCount = records.count()

//transformation - datasets are immutable!
val prs = records.where(records("type") === "PullRequestEvent")
//action - job 2
val pullRequestEventCount = prs.count()
```


Job 0 - Stages, tasks, partitions

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at HelloSparkWorld.scala:13 count at HelloSparkWorld.scala:13	2019/04/30 22:53:57	4 s	2/2	4/4
0	json at HelloSparkWorld.scala:11 json at HelloSparkWorld.scala:11	2019/04/30 22:53:53	3 s	1/1	3/3

Job 0 - Stage 0, tasks, partitions

Stage Id ▼	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
0	json at HelloSparkWorld.scala:11 +details	2019/04/30 22:53:53	3 s	3/3	46.6 MB			

Input partitions - splittable file?



Job 1 - count

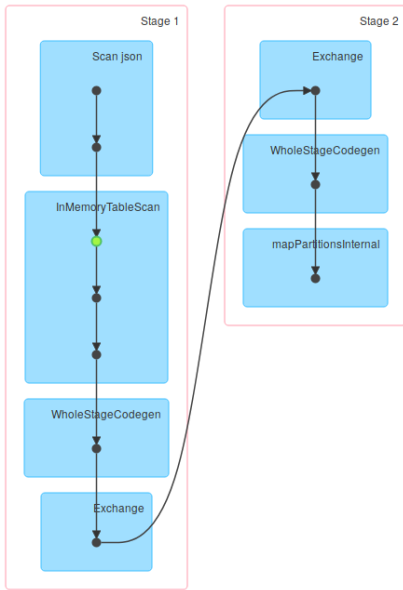
Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
1	count at HelloSparkWorld.scala:13 count at HelloSparkWorld.scala:13	2019/04/30 22:53:57	4 s	2/2	4/4
0	json at HelloSparkWorld.scala:11 json at HelloSparkWorld.scala:11	2019/04/30 22:53:53	3 s	1/1	3/3

Job 1 - count Stages 1 and 2

Stage Id ▼	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
2	count at HelloSparkWorld.scala:13 +details	2019/04/30 22:54:02	69 ms	1/1			177.0 B	
1	count at HelloSparkWorld.scala:13 +details	2019/04/30 22:53:57	4 s	3/3	46.6 MB			177.0 B

Job 1 - Stages 1 and 2 DAG

▼ DAG Visualization



RDDs - Not deprecated!

```
object RddProcessor {  
  
  val DefaultEventInputUrl = "file:///datasets/github/data"  
  
  def process(sc: SparkContext, inputUrl: String): (Long, Long) = {  
    val records: RDD[String] = sc.textFile(inputUrl)  
    println(s"We have a total of ${records.partitions.size} partitions.")  
    val total = records.count()  
    val prs = records.filter(r => r.contains("PullRequestEvent"))  
    val totalPrs = prs.count()  
    (total, totalPrs)  
  }  
  
  def main(args: Array[String]): Unit = {  
    val spark = SparkSession.builder().  
      appName(name = "RddProcessor").  
      getOrCreate()  
  
    val inputUrl = if (args.size > 0) {  
      args(0)  
    } else {  
      DefaultEventInputUrl  
    }  
    process(spark.sparkContext, inputUrl)  
    spark.stop()  
  }  
}
```

And now for something completely different: Colon Cancer



- Screening saves lives!
 - Colonoscopy - talk to your doc
 - Dave Barry: A journey into my colon — and yours
- Colorectal Cancer Alliance

Questions?



- medale@asymmetrik.com
- Infrequent blog/past presentations <http://uebercomputing.com/>